

# Supplementary Information: Experiment-Driven Bayesian Optimization of the Impact Response of Multi-Material 3D-Printed Tensegrity-Inspired Structures

**Draft note (remove before submission):** as in the main text, any value printed in this violet color is a synthetic placeholder rather than a measurement. As of 2026-09-15 this Supplementary Information contains no violet values: the former placeholder batch-2 print key was replaced by the logged campaign record, which now covers batches 2 to 4.

This Supplementary Information collects the detailed design-, fabrication-, test-, and analysis-development records supporting the main manuscript. Unlike the main text, the SI references the project’s development history directly, including specific pull requests (PRs) and issues in the project repository (<https://github.com/vertical-cloud-lab/tensegrity-optimization>).

## S1 Joint-Design Study

The strut-to-cable junction was developed and ranked through a five-design OpenSCAD study (anchor-bulb, dovetail, TPU-sleeve overmold, eyelet-loop, and TPU-rebar variants). The working prototype uses a dovetail joint (Design B) with an anchor-bulb backup (Design A) for sensitivity studies; a captive TPU core is routed inside the PLA shell so that the soft tendon is protected from layer-line failure at the transition.

Full results, renders, and the Phase-3/Phase-4 design reviews are documented in [PR #39](#) (joint design) and [PR #35](#) (CAD parameterization and print history).

## S2 Support Strategy for Soft Members

Near-vertical TPU tendons are otherwise unsupported during printing. For the printed campaign batches, automatic support generation was disabled and PLA edge supports were painted manually in Bambu Studio along each tendon, with the support-to-object XY clearance tightened from 0.35 mm to 0.2 mm during tuning on the S0 reference prism ([issue #98](#)). Support removal is the leading source of the logged cosmetic defects (surface stringing; on two seed articles, TPU strands pulled off bottom tendons), so defects are recorded per article in the print key (Section S3).

An automated alternative was developed in parallel: a tensegrity-specific Bambu Studio recipe (support threshold angle lowered from 40° to 10°, support material matched to the rigid extruder) combined with manually generated narrowing pillars that ray-cast against the printable mesh underside. The narrowing-pillar workflow is documented in [PR #66](#); the H2D multi-part material assignment fix is in [PR #64](#).




Figure S1: Five candidate strut-to-cable joint geometries from the OpenSCAD joint study (left to right: A anchor-bulb, B dovetail, C TPU-sleeve overmold, D eyelet-loop, E TPU-rebar; top row isometric, bottom row section views). The study renders use the PETG (orange) + TPU (cyan) material pair under which the joint family was first developed; the working prototype adapts the selected dovetail (B) with an anchor-bulb (A) backup to the PLA/TPU pair of the main text, with the captive TPU core routed inside the PLA shell.

### S3 As-Printed Campaign Record (Print Key)

Table S1 is the print key for the seed campaign: every article's print identifier, its design, its weighed mass, and its logged defects. The authoritative machine-readable key, together with the as-printed slicer project (one specimen per named build plate), lives at [PR #102](#) (`bo/t3-prism-bo-batch-print-key.csv`). Two entries deserve care. Design 08 was printed three times; the plate label in the archived slicer project and a contemporaneous issue comment disagree on which article is the official one (*dea4ls* versus *bag26v*), and pending confirmation the analysis uses *bag26v*. The tri-print of design 08 also supplies the only print-to-print repeatability number to date: a mass standard deviation of 0.457 g (three prints of a single geometry, so not a general process repeatability estimate). Separately, one tested article (*amdjwm*) has no confirmed design mapping and is excluded from the surrogate fit until identified.

Table S2 is the corresponding logged key for the three model-recommended batches. Every batch-2 and batch-3 article printed at about 11% dry-box relative humidity and every batch-4 article at about 7%. The batch-3 plate was printed twice; the reprint articles (2dran1 to 2dran9, weighed 19.48 to 20.32 g, 0.27 g heavier than the first print article for article on average, all logged defect-free) supplied the second test pass of the main text. The lab record confirms each 2dran $N$  as the re-print of drran $N$  at the same Ax trial, so the two passes pair article for article (tabulated in the main-text Results), and the nine reprint articles subsequently entered the surrogate fit as distinct completed trials at their own weighed masses and per-session standard errors. The authoritative machine-readable keys, with plate positions and the mapping evidence per row, are the `t3-prism-bo-round1/round3/round4-print-key.csv` files archived with the manuscript data.

Table S1: Print key for the seed campaign (nine Sobol designs plus the S0 reference). Masses weighed on a 0.01 g scale. TPU fed from a heated dry box near 8% relative humidity where logged. The first five articles (printed before the TPU nozzle fix) used the standard 0.4 mm TPU nozzle; the rest used the high-flow 0.6 mm nozzle.

| Design  | Article | Mass (g) | Logged defects                                   |
|---------|---------|----------|--------------------------------------------------|
| Spec 00 | 9hhbkp  | 21.62    | none                                             |
| Spec 01 | 6lhxfy  | 18.50    | slight strut misprint; divot on one top tendon   |
| Spec 02 | autv5r  | 22.04    | minor tendon misalignment near top               |
| Spec 03 | ebdna8  | 19.61    | two strands off bottom tendons (support removal) |
| Spec 04 | nvxsr   | 20.66    | misalignment ring after printer pause            |
| Spec 05 | 6nheas  | 21.73    | two strands off bottom tendons (support-related) |
| Spec 06 | lzm8rv  | 18.50    | stringing along diagonal tendons                 |
| Spec 07 | ajhby6  | 20.53    | filament line pulled from bottom struts          |
| Spec 08 | dea4ls  | 22.29    | stringing; residual PLA on tendon                |
| Spec 08 | bag26v  | 21.42    | some stringing on tendon                         |
| Spec 08 | ghmj4y  | 22.10    | inconsistent tendon cross-section near top       |
| S0 ref. | bpx68c  | 20.23    | none major                                       |

## S4 Drop-Tower Protocol and Rig Characterization

The campaign fixture is a Lansmont Model 23 shock test system with a Test Partner 4 data-acquisition console. The development history behind the protocol in the main text is summarized here.

**Protocol.** 60 in drop height; specimen upright on the carriage plate; first two drops of each session discarded as seating warm-up. Session length by batch: seed sessions contribute 101 valid captures per article (103 for autv5r); the batch-2 record lists 21 to 22 valid captures per article; batch-3 and batch-4 sessions recorded 20 drops per article. Exact per-session counts live in the archived per-drop records, whose two generations differ on whether the count is quoted before or after the warm-up discard. The shortening is justified in the main text (Section 3.3): the statistical minimum from the pre-campaign variance study is five stabilized drops, a truncation replay of the seed sessions preserves the full-session transmissibility ranking at twenty drops, and the batch-3 second pass showed the error budget is dominated by seating and print variation rather than session length. Three slow-motion video clips (959 fps) per specimen (from batch 4 filmed at drops 1, 10, and 20 so clips pair deterministically with captures); roughly 42 s per drop, so about ninety minutes per seed article and fifteen per later-batch article. Capture settings: 1.25 MHz sample rate, 100 ms record, 2 ms pre-trigger, 150 G trigger on the base channel, AC-coupled ICP conditioning. The console settings were re-verified digit-for-digit against an archived screenshot after a settings-loss event, with data continuity confirmed to within about 1%. Per the review instruction of 2026-09-02, the per-session capture counts above should be re-verified by a person against the committed per-drop records before submission.

**Impact surface.** A blinded  $3 \times 3$  crossover experiment over three candidate polyurethane mat arrangements (single 0.25 in Shore 70 OO; single 0.5 in Shore 60 OO; stacked) selected the single 0.5 in Shore 60 OO sheet ([issue #88](#), [PR #86](#)). Unlike the earlier felt stack, the mat does not measurably compact over hundred-drop sessions. Historical note: sessions recorded as “5 felts” in early protocol documents were physically four felt sheets plus one cardboard layer.

Table S2: Print key for the model-recommended batches (first prints; the batch-3 reprint set is summarized in the text). Masses weighed on a 0.01 g scale.

| Batch | Trial | Article | Mass (g) | Logged defects                                         |
|-------|-------|---------|----------|--------------------------------------------------------|
| 2     | 10    | r2d2c9  | 17.91    | none                                                   |
| 2     | 11    | r2d2c4  | 19.33    | string on one top tendon; partial TPU detachment       |
| 2     | 12    | r2d2c2  | 17.96    | none                                                   |
| 2     | 13    | r2d2c1  | 19.38    | detached TPU string, two top tendons; odd foot support |
| 2     | 14    | r2d2c6  | 19.11    | none                                                   |
| 2     | 15    | r2d2c7  | 18.51    | none                                                   |
| 2     | 16    | r2d2c3  | 18.63    | none                                                   |
| 2     | 17    | r2d2c8  | 19.98    | full-span strings, two bottom tendons                  |
| 2     | 18    | r2d2c5  | 23.47    | none                                                   |
| 3     | 28    | drran5  | 19.88    | tiny bubbles, diagonal tendons                         |
| 3     | 29    | drran4  | 19.92    | none                                                   |
| 3     | 30    | drran9  | 19.34    | none                                                   |
| 3     | 31    | drran3  | 19.92    | tiny bubbles, diagonal tendons                         |
| 3     | 32    | drran7  | 20.06    | tiny bubbles, diagonal tendons                         |
| 3     | 33    | drran2  | 19.33    | none                                                   |
| 3     | 34    | drran6  | 19.24    | tiny bubbles, diagonal tendons                         |
| 3     | 35    | drran8  | 19.32    | none                                                   |
| 3     | 36    | drran1  | 19.31    | none                                                   |
| 4     | 37    | corny7  | 19.62    | none                                                   |
| 4     | 38    | corny2  | 20.21    | none                                                   |
| 4     | 39    | corny8  | 19.64    | none                                                   |
| 4     | 40    | corny5  | 20.58    | none                                                   |
| 4     | 41    | corny6  | 19.45    | none                                                   |
| 4     | 42    | corny1  | 20.20    | TPU mixed into one PLA strut layer                     |
| 4     | 43    | corny9  | 19.46    | none                                                   |
| 4     | 44    | corny4  | 20.56    | small pores, diagonal tendons                          |
| 4     | 45    | corny3  | 20.15    | none                                                   |

**Sensors and calibration.** Base input: single-axis piezoelectric accelerometer on the carriage plate (1.059 mV/G, full scale 9443 G); top response: Dytran 3133A4 triaxial accelerometer (X/Y/Z 0.690/0.667/0.734 mV/G) seated in a printed pocket at the top vertex, mounted per ISO 5348. An earlier single-axis channel was retired after range clipping. The two base-capable channels were cross-calibrated in place over fifteen co-located drops: ratio  $0.953 \pm 0.001$  at CFC-180 (per-drop  $0.952 \pm 0.004$ ), stable across drop heights ([PR #74](#)).

**Filtering correction.** The SAE J211/1 channel-frequency-class filter used for all reported numbers is implemented from the standard’s published coefficients. An earlier in-house implementation passed the two-pass (forward-backward) corner frequency to a single-pass design, which narrows every channel class by roughly 20% (an effective CFC-144 where CFC-180 was intended); the discrepancy was identified during the analysis audit of [issue #94](#) and corrected before the campaign numbers were produced.

**Processing pipeline.** For each capture, every channel is calibrated and CFC-filtered independently; the top-vertex magnitude  $\|\mathbf{a}_{\text{top}}(t)\|$  is formed from the three filtered triaxial components, and the two peaks entering  $t_{180}$  are searched within a short window (approximately 1.5 ms) of first contact, so the peaks need not be simultaneous. The baseline is estimated from the record tail (at healthy tower speed the mat-contact foot begins before the 2 ms pre-trigger window).  $\Delta v$  is the integral of the baselined base channel over the impact pulse and doubles as the rig-health gate. Drops are discarded per the warm-up rule. Table S3 collects the operational analysis parameters in one place: values known from the protocol are stated numerically, and the remaining thresholds are to be transcribed verbatim from the version-tagged campaign analysis code at its archival release (a declared submission gate), rather than paraphrased here at the risk of diverging from the code of record.

Table S3: Operational analysis parameters for the drop-tower pipeline. “Code of record” marks values to be transcribed from the version-tagged analysis release before submission.

| Parameter                                 | Value                                                |
|-------------------------------------------|------------------------------------------------------|
| Drop height                               | 60 in (1.524 m)                                      |
| Sample rate / record / pre-trigger        | 1.25 MHz / 100 ms / 2 ms                             |
| Trigger                                   | 150 G on the base channel                            |
| Reporting filter class                    | CFC-180 (SAE J211/1), two-pass corner corrected      |
| Diagnostic filter class                   | CFC-1000                                             |
| Peak-search window                        | $\approx 1.5$ ms from first contact                  |
| Baseline estimate                         | record tail                                          |
| Warm-up exclusion                         | first 2 drops per session                            |
| Captures per session                      | 101 seed (103 autv5r); 21-22 batch 2; 20 batches 3-4 |
| Rig-health check                          | $\Delta v$ from baselined base channel               |
| Ringdown fit gate                         | $r^2 \geq 0.85$ (exponential fit)                    |
| Second-contact threshold and window       | code of record                                       |
| Per-capture exclusion rules               | code of record                                       |
| Stabilized-mean convention                | code of record                                       |
| Filter initialization / boundary handling | code of record                                       |

**Rig health and repair history.** Rig health is checked against the impact velocity change  $\Delta v$  recovered from the base channel; over an uninterrupted 100-drop session the  $\Delta v$  CV was 0.54% with a flat drift slope. Within a session,  $\Delta v$  and the CFC-180 input typically ease down by 1 to 4% as the polyurethane mat warms, and recover fully after a rest of minutes; after a mat rest or re-seat, the first 5 to 10 drops carry the settling transient, which is why the warm-up discard exists. The arrest velocity has visibly stabilized as cumulative drop count grew, which is the observation behind keeping the  $\Delta v$  record even though it is no longer listed among the paper’s contributions. During commissioning, two guide-rod locking pins failed and a defeated pin-park safety interlock was discovered and recommissioned; the pin-related friction produced a measurable  $\Delta v$  deficit that resolved after repair ([issue #92](#)). More than 4,400 instrumented drops across roughly 80 printed specimens have been performed over about four months of commissioning, characterization, and campaign sessions (tallied through [issue #101](#) plus the batch records); the campaign batches themselves contribute about 1,500 stabilized captures including the batch-3 second pass.

**Seating sensitivity, the between-seat rule, and the seat gauge.** The batch-3 second pass (main text, Section 4.2) established that a session can be internally stable and still misrepresent

its design: the 1.251 outlier session passed every within-session check and did not reproduce. Two standing conventions followed. First, the *between-seat replication rule*: any single-seating value that would drive a decision requires confirmation on an independent re-seat. Second, an advisory *seat gauge*: the broadband ratio of CFC-1000 to CFC-180 transmissibility reads 1.00 to 1.07 in healthy sessions, and every large same-design shift observed so far carried a gauge of at least 1.15 in one of its seatings, so gauge-hot sessions are treated as re-seat candidates. Both conventions, with the session-by-session drift-watch record, live in the drop-test datasets on [PR #86](#).

**Objective provenance and adversarial review.** The objective stack evolved through a measurability review ([PR #97](#)): per-specimen crush energy in joules is not recoverable from this rig (no falling-mass or displacement channel, and the specimens stay elastic), which is why the campaign responses are the filtered peak-acceleration ratio and a provisional mass-weighted rebound-timing score, with quasi-static crush metrics deferred to a load-frame campaign. A pre-registered adversarial review of the objective choices (Edison ANALYSIS task [3e398131](#), archived in the repository) recommended accounting for article-level rather than per-drop noise, treating the rebound fraction as a candidate constraint pending a restrained-versus-unrestrained validation with synchronized video, and reserving print capacity for replicate articles. The campaign’s realized responses are recorded in the main text: the batch-3 plate was printed and tested twice, which measured the design-level shift distribution directly and produced the between-seat replication rule; the rebound fraction remains a candidate constraint; and the surrogate continues to ingest per-session standard errors, with the replication evidence carried as an interpretation bound rather than as an inflated noise term. The reserved-replicate recommendation now has a planned instrument: the replicate rank-stability study of the main-text Methods will re-print the best-, mid-, and worst-predicted in-sample designs in nominal triplicate and test whether the predicted ordering survives print-to-print and seating variation.

## S5 Printed-Mass Model

The seed batch was projected onto constant *solid* CAD mass ( $m^* = 30.95$  g), yet the articles weigh 18.50 g to 22.29 g, because PLA struts print at partial infill (15% grid, two walls) while thin TPU cables print near solid, and the PLA-to-TPU volume split swings widely across the design box. A two-stage calibrated model closes this gap ([PR #102](#), `bo/t3-prism_mass_model.py`): analytic member volumes are corrected to rendered-solid grams (one factor per material), and rendered-solid grams are mapped to weighed grams through a wall-plus-infill term in the as-printed strut diameter (fitted infill 23.6%, wall 1.18 mm, TPU factor 0.744). The model’s residual standard deviation is 0.378 g. For context, the three available prints of design 08 have a sample standard deviation of 0.457 g; three prints of one geometry are insufficient to establish a general process repeatability limit. Inverting the model inside the projection lets a batch hold *printed* grams constant with no slicer in the loop, and batches 3 and 4 did so at the 20.23 g target (the reference article’s weight). Outcomes: batch mass CV fell from 5.5% (tested seed articles) and 8.7% (batch 2, still solid-mass-projected) to 1.7% (batch 3) and 2.3% (batch 4). The weighings sit a few tenths of a gram below target as a per-session flow offset (−0.52, −0.25, and −0.26 g for the batch-3 first print, its reprint, and batch 4), while the model’s infill terms passed their first test: regressing the print-averaged mass residuals of the eighteen batch-3 articles on the two commanded infills gives slopes indistinguishable from zero over the realized 13 to 34% spread. A session-term recalibration with the accumulated weighings is queued before the next batch’s projection. Two batch-2 recommendation sets exist in the record and must not be conflated: the batch actually sliced and sent to the

printer was generated before this correction, under the seed batch’s solid-mass projection; a subsequently regenerated set applies the constant-printed-mass projection and seeded the batch-3 tooling. Both are archived with the manuscript data (`round2-as-printed-solid-mass-projection.csv`, `round3-candidate-constant-printed-mass.csv`).

## S6 Metal-Analog Validation Plan

A planned future validation campaign rebuilds a subset of the  $T_3$ -prism designs as metal analogs (hollow aluminum rods for the compression members and stainless-steel threaded cables for the tension members) using the same  $R$ ,  $H$ ,  $\theta$ , and member-diameter parameterization as the printed PLA/TPU specimens. To probe transfer across material systems rather than only confirming the optimum, the selected designs deliberately span the predicted performance range: representatives the surrogate predicts to be worst-, mid-, and best-performing are each built in both material systems and characterized under the identical drop protocol of Section S4. Rank agreement between the measured PLA/TPU ranking and the hollow-aluminum/stainless-steel ranking on the primary campaign metric  $t_{180}$  (to be reported in the main-text metal-analog table and figure, with the provisional rebound score and, later, quasi-static SEA as secondary rankings) would be preliminary evidence consistent with transfer beyond the as-printed polymers; disagreement would not isolate which material or joint difference caused the change.

*Placeholder: assembled metal-analog specimens (hollow aluminum rods + stainless-steel threaded cables) beside their printed PLA/TPU equivalents for the worst/mid/best predicted tiers (photographs pending fabrication).*

Figure S2: Metal-analog validation specimens alongside their printed PLA/TPU equivalents for each predicted performance tier.